


Probe specificity breaks on authorized actions



Separate evaluation panels; do not pool counts. Earlier endpoints use their own calibration.

Bars describe observed cases, not deployment rates. Few task families and zero observed errors limit generalization.